

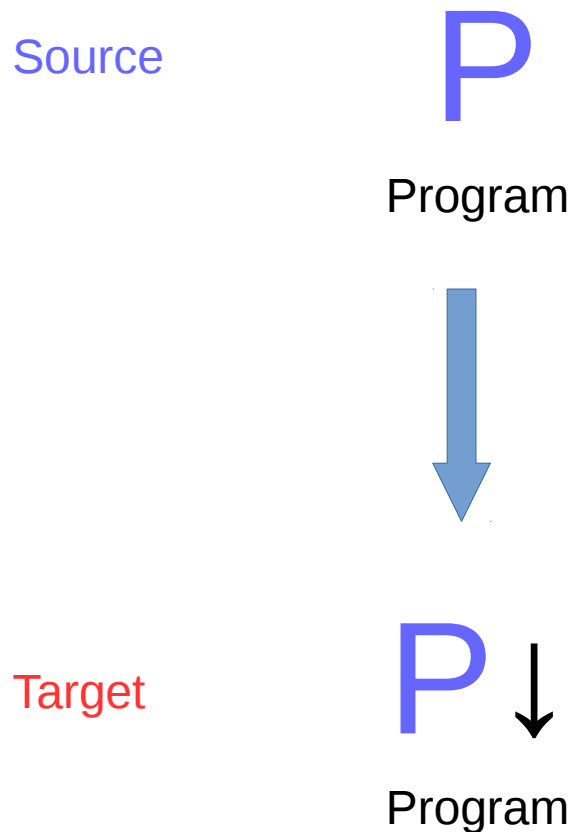
How to *define* compiler security?

(A property-centric view)

Ongoing work with:

Catalin Hritcu, Marco Patrignani, Jeremy Thibault,
Carmin Abate, Roberto Blanco + others

An overview




Correctness: Preserve P 's properties

If P has a property, then $P \downarrow$ has the property.

Security: Preserve P 's properties despite adversaries


If P has a property despite all source adversaries, then $P \downarrow$ has the property despite all target adversaries.

Observable behavior as traces



- Trace: Sequence of actions representing behavior observed by environment.
?0 !0 ?5 !5 ?2 !7 ?3 !10 ...
?1 !1 ?2 !2 ?3 !6 ?7 !42 ?8 !(Halt)
- $\text{Tr}(\textcolor{blue}{P})$ = Set of all traces of $\textcolor{blue}{P}$
- $\text{Tr}(\textcolor{blue}{P})$ describes $\textcolor{blue}{P}$'s observable behavior
- Does not represent hidden state and unobservable ways of interacting with $\textcolor{blue}{P}$

Properties as sets of traces



- Property/specification π : Set of good traces

$$\pi = \{ ?a_0 !b_0 ?a_1 !b_1 \dots \mid b_i \geq 0 \}$$

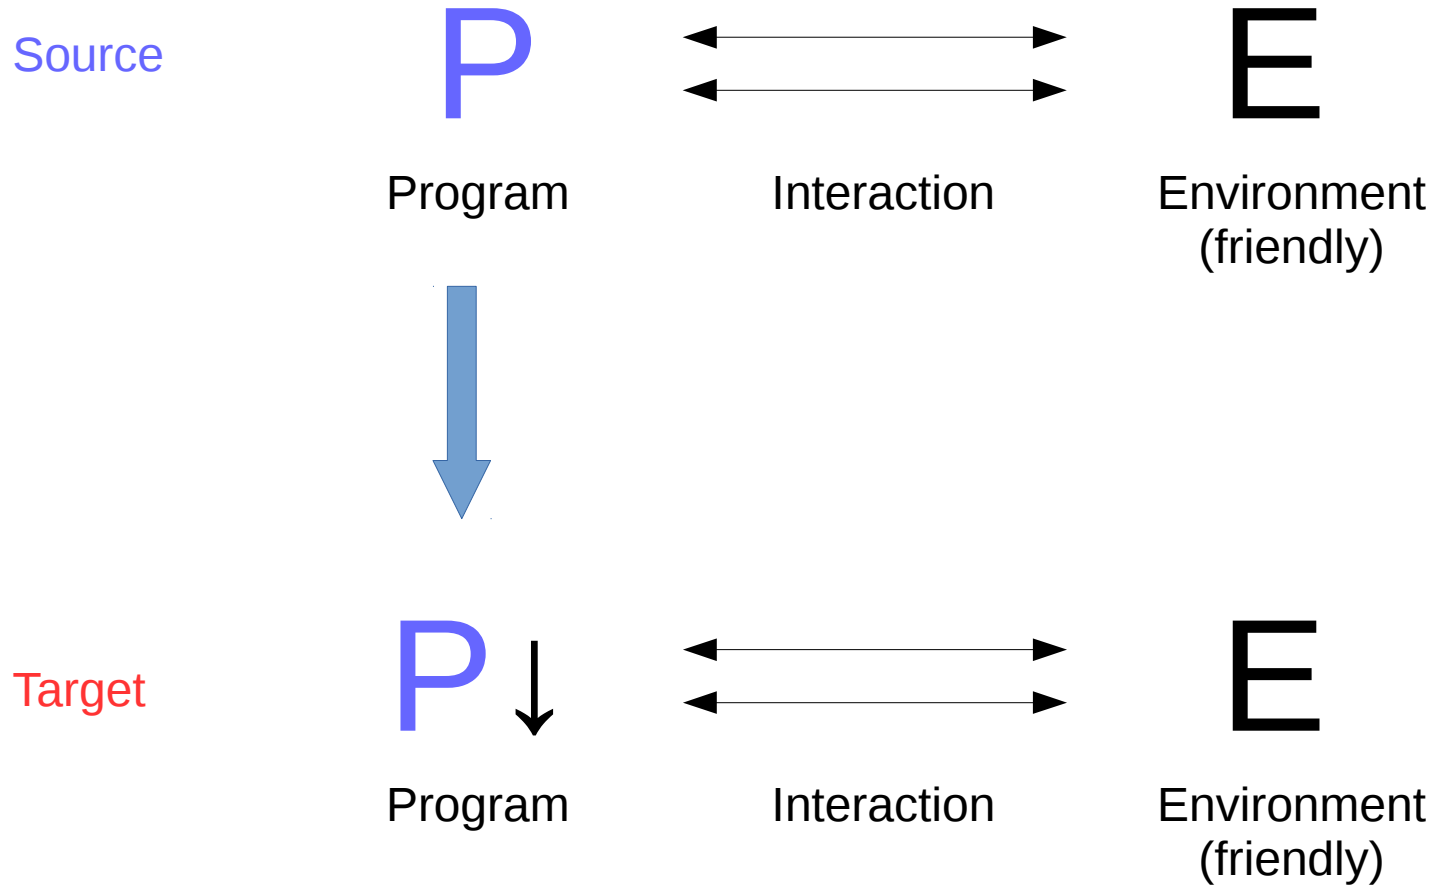
(Non-negative outputs only)

$$?0 !0 ?5 !5 ?2 !7 \dots \in \pi$$

$$?0 !0 ?5 !(-5) \dots \notin \pi$$

- P has property π iff $\text{Tr}(P) \subseteq \pi$

Compilation



Compiler correctness as property preservation

A compiler \downarrow is correct if it *preserves all properties*:

$$\forall \pi. \forall P. (P \text{ has } \pi) \Rightarrow (P \downarrow \text{ has } \pi)$$

Key idea

If we checked (manually or by verification) that the source program has π , then there is no need to re-verify the compiled program.

Compiler correctness: Equivalent characterization

- A compiler \downarrow is correct if it *preserves all properties*

$$\forall \pi. \forall P. (P \text{ has } \pi) \Rightarrow (P \downarrow \text{ has } \pi)$$

- Equivalently,

$$\forall \pi. \forall P. (\text{Tr}(P) \subseteq \pi) \Rightarrow (\text{Tr}(P \downarrow) \subseteq \pi)$$

- Equivalently,

$$\forall P. \text{Tr}(P \downarrow) \subseteq \text{Tr}(P)$$

This last statement is called “refinement”

Theorem

A compiler *preserves all properties* iff it *refines* behaviors (only reduces traces)

Compiler correctness: Equivalent characterization

- A compiler \downarrow is correct if it *preserves all properties*

$$\forall \pi. \forall P. (P \text{ has } \pi) \Rightarrow (P \downarrow \text{ has } \pi)$$

A property-based
characterization

```
graph TD; A["A property-based characterization"] --> B["A property-free characterization"]; B --> C["Amenable to formal proof"]; C --> D["Equivalently, Tr(P↓) ⊆ Tr(P)"]; D --> E["Equivalently, Tr(P↓) ⊆ Tr(P)"];
```

A property-free
characterization


- Equivalently,

$$\forall P. \text{Tr}(P \downarrow) \subseteq \text{Tr}(P)$$

Amenable to
formal proof


Security!

Behavior as traces, with adversaries




- **A** may be:
 - A library co-linked with **P** (corrupts internal state of **P**)
 - A malicious network client (provides exploit inputs)
- **A**'s interactions use a possibly different/broader API than **E**'s interactions
- **A**'s interactions are *not* what we want to specify, so they are not represented on the traces (gives compiler leeway)

Behavior as traces, with adversaries



- Example:
 - **P** accumulates inputs in a counter, outputs counter value at each step. Ideally,
?0 !0 ?5 !5 ?2 !7 ?3 !10 ...
 - **A** is a library that **P** uses to log on the side (for debugging)
 - But, due to a buffer overflow bug, **A** also occasionally *corrupts* **P**'s internal counter, breaking **P**'s expected behavior.
- The $A \leftrightarrow P$ and $P \leftrightarrow E$ interactions are different.
- The $A \leftrightarrow P$ interaction is internal, not part of traces or properties.

Properties, with adversaries




- Properties/specifications π same as without **A**.
- But notion of program having property changes:

P *robustly has* π iff $\forall \mathbf{A}. \text{Tr}(\mathbf{A} || \mathbf{P}) \subseteq \pi$


(cf. robust safety)

Compilation, with adversaries



May be more powerful than source adversaries

A secure compiler ensures that adversary's increased power does not break expected $P \leftrightarrow E$ properties



Compiler security as robust property preservation

A compiler \downarrow is secure if it *preserves all properties **robustly***:

$$\forall \pi. \forall P. (P \text{ robustly has } \pi) \Rightarrow (P \downarrow \text{ robustly has } \pi)$$

$$\forall \pi. \forall P. (\forall A. \text{Tr}(A \parallel P) \subseteq \pi) \Rightarrow (\forall A'. \text{Tr}(A' \parallel (P \downarrow)) \subseteq \pi)$$

For all π , the compiled program is no more vulnerable to attacks on π than the source program. Hence, verifying the source (in source semantics) is enough to guarantee π even against low-level attacks.

Compiler security: Equivalent characterization

- A compiler \downarrow is secure if it *preserves all properties **robustly***

$$\forall \pi. \forall P. (P \text{ robustly has } \pi) \Rightarrow (P \downarrow \text{ robustly has } \pi)$$

- Equivalently,

$$\forall \pi. \forall P. (\forall A. \text{Tr}(A \parallel P) \subseteq \pi) \Rightarrow (\forall A'. \text{Tr}(A' \parallel (P \downarrow)) \subseteq \pi)$$

- Equivalently,

$$\forall A' P t. (t \in \text{Tr}(A' \parallel (P \downarrow))) \Rightarrow \exists A. t \in \text{Tr}(A \parallel P)$$

We call this last statement “robust refinement”

Theorem

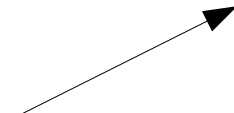
A compiler *preserves all properties robustly* iff it *refines behaviors robustly*.

Compiler security: Equivalent characterization

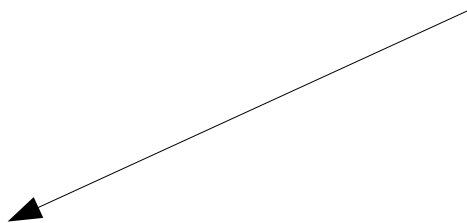
- A compiler \downarrow is secure if it *preserves all properties **robustly***

$$\forall \pi. \forall P. (P \text{ robustly has } \pi) \Rightarrow (P \downarrow \text{ robustly has } \pi)$$

A property-based
characterization



A property-free
characterization



- Equivalently,

$$\forall A' P t. (t \in \text{Tr}(A' \parallel (P \downarrow))) \Rightarrow \exists A. t \in \text{Tr}(A \parallel P)$$

Related to “back-translation”



Amenable to formal proof

Vertically composable

Compiler security, parametrized

- Does security coincide with preserving all properties robustly?
 - Not necessarily: It may require less (or more ... later)
- A compiler \downarrow is secure wrt a class C of properties if it *robustly preserves all properties in C*
$$\forall \pi \in C. \forall P. (P \text{ robustly has } \pi) \Rightarrow (P \downarrow \text{ robustly has } \pi)$$
- C may be “safety properties” or “liveness properties”
- Standard notions from literature

Safety properties

- Intuitively, a property that specifies “nothing bad will happen”
- “Bad” = Set M of finite trace prefixes
- Corresponding safety property:

$$\pi_M = \{t \mid \forall m \in M. \text{not } (m \leq t)\}$$

(Traces that don't extend anything from M)

- Example:

$$M = \{t \mid t \text{ finite and } t \text{ ends in a negative output}\}$$

$$\pi_M = \{t \mid t \text{ does } \textit{not} \text{ have a negative output}\}$$

Compiler security for safety

- A compiler \downarrow is secure for safety if it *preserves all safety properties robustly*

$$\forall \pi \in \text{Safety}. \forall P. (P \text{ robustly has } \pi) \\ \Rightarrow (P \downarrow \text{ robustly has } \pi)$$

- Equivalently,

$$\forall A' P. t. (t \in \text{Fin}(A' \parallel (P \downarrow))) \Rightarrow \exists A. t \in \text{Fin}(A \parallel P)$$

Theorem

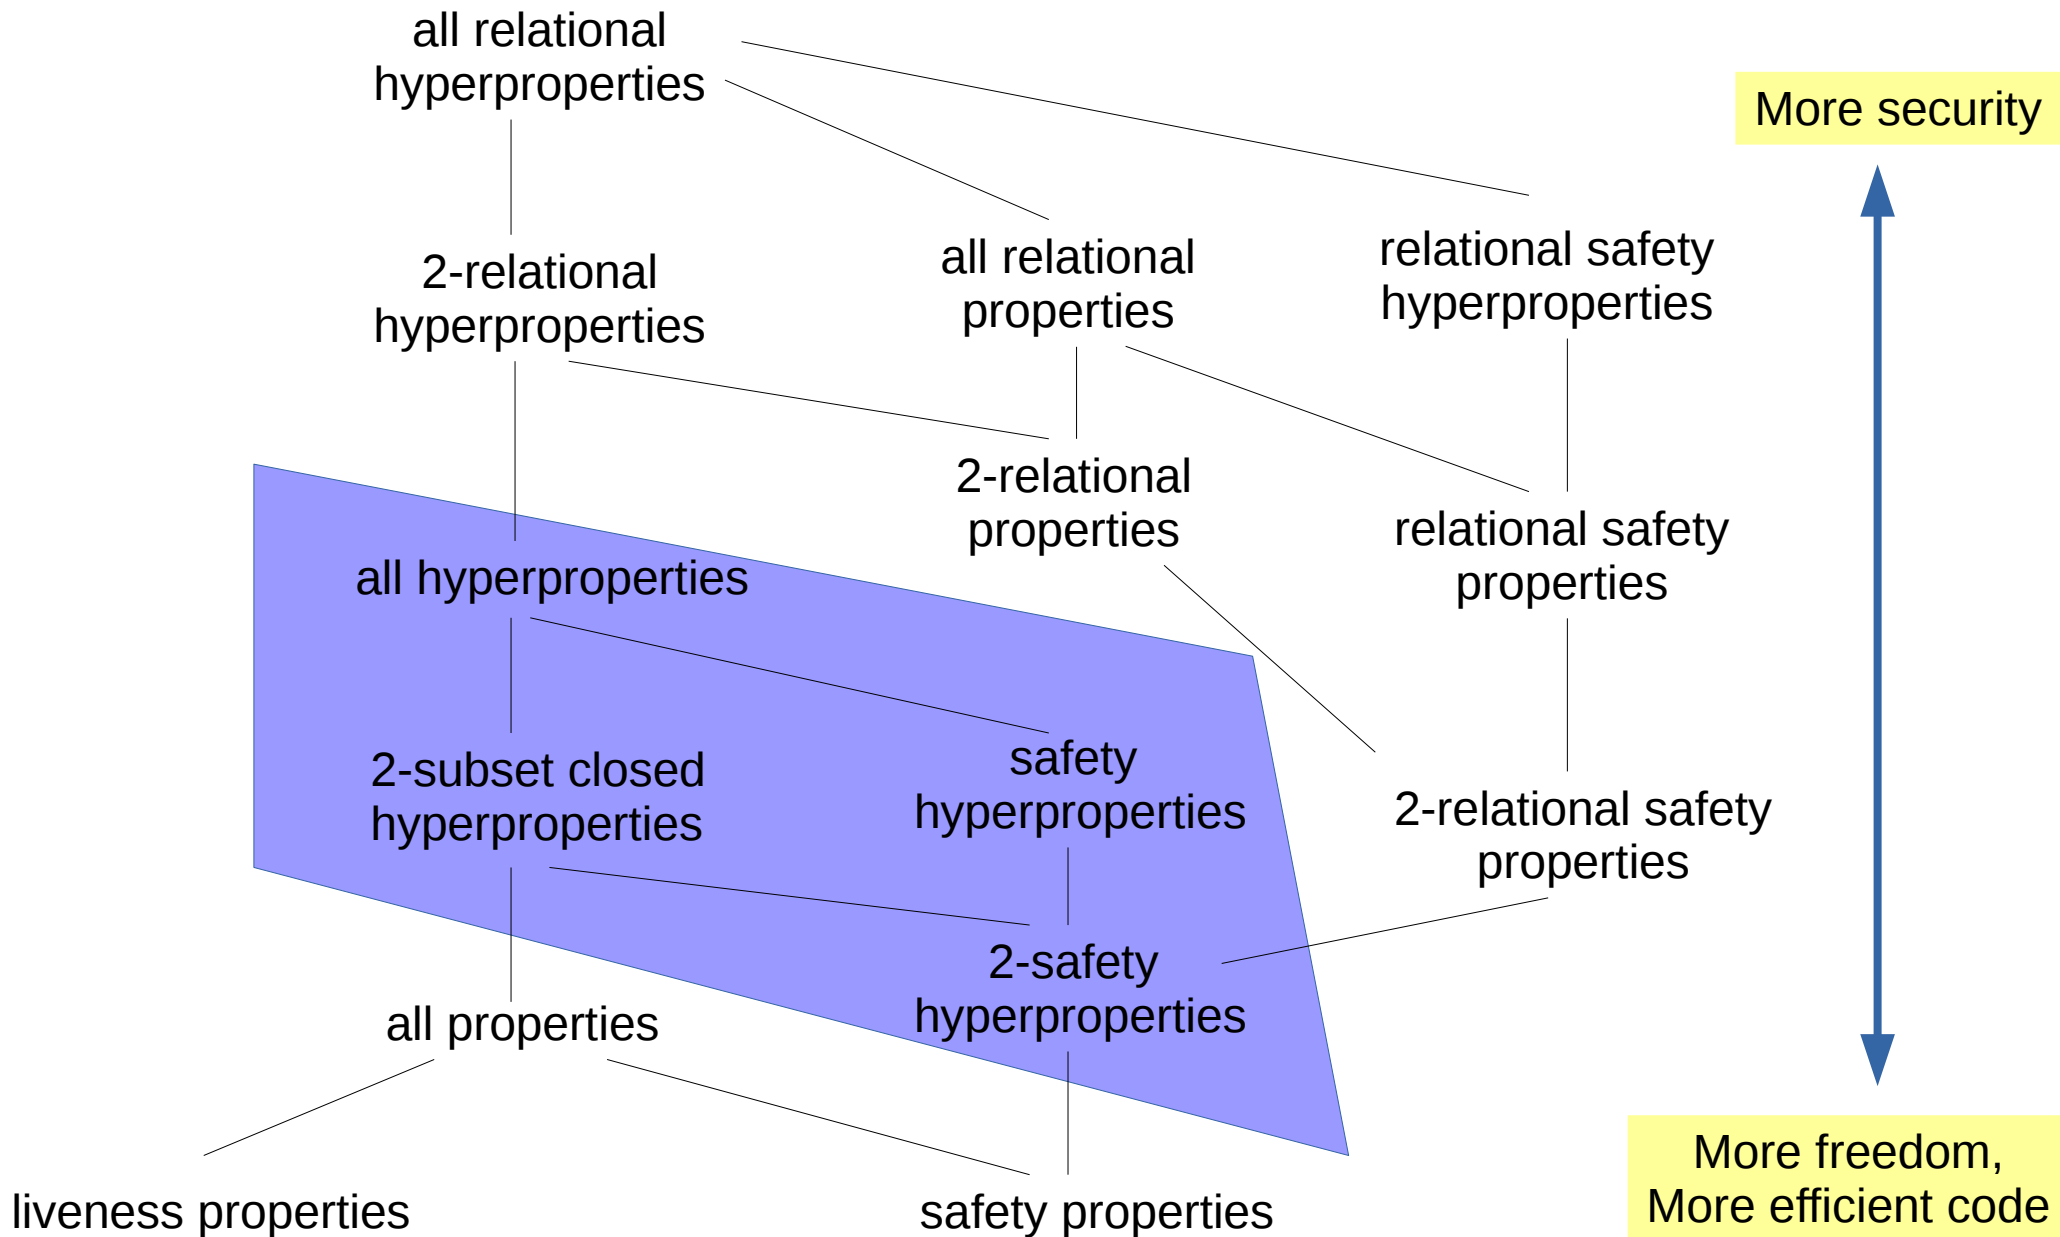
A compiler *preserves all safety properties robustly* iff it *refines finite behaviors robustly*.

Defining compiler security: General principle

- Define compiler security as preservation of some robust security class X .
- Obtain an equivalent, X -free characterization that is proof-amenable.
- Surprisingly, can be done for many different X .

Summary of results (abridged)

A compiler is secure for X if it robustly preserves all X, where X is:



2-safety hyperproperties (2-SHP)

- A 2-safety hyperproperty H is characterized by a set of *pairs* of bad trace prefixes, T as follows:
 P has H iff there is no pair in T such that P has both traces of the pair.
- Example: Program takes two inputs and produces one output:
 - 2-SHP: “Output must be independent of first input”
 - $T = \{(?i_1 \ ?i_2 \ !o_1, \ ?i_1' \ ?i_2 \ !o_1') \mid o_1 \neq o_1'\}$
 - Confidentiality property

Compiler security for 2-SHP

- A compiler \downarrow is secure for 2-SHP if it *preserves all 2-SHPs robustly*

$$\forall H \in 2\text{-SHP}. \forall P. (P \text{ robustly has } H) \\ \Rightarrow (P_{\downarrow} \text{ robustly has } H)$$

- Equivalently,

$$\forall A' P t_1 t_2. (t_1, t_2 \in \text{Fin}(A' \parallel (P_{\downarrow}))) \\ \Rightarrow \exists A. t_1, t_2 \in \text{Fin}(A \parallel P)$$

Compare to compiler security for safety:

$$\forall A' P t. (t \in \text{Fin}(A' \parallel (P_{\downarrow}))) \Rightarrow \exists A. t \in \text{Fin}(A \parallel P)$$

More examples

- Security for all hyperproperties

$$\forall A' P. \exists A. \text{Tr}(A' \parallel (P_{\downarrow})) = \text{Tr}(A \parallel P)$$

- Security for all relational hyperproperties

$$\forall A'. \exists A. \forall P. \text{Tr}(A' \parallel (P_{\downarrow})) = \text{Tr}(A \parallel P)$$

Current status

One compiler security criterion for each point in the partial order. Equivalent, back-translation-like characterization.

What must a compiler do to attain each point?

- Safety: Protect integrity of P_{\downarrow} 's private state
- 2-safety: Above + protect confidentiality of P_{\downarrow} 's private state
- Relational properties: Above + protect code identity

Compilers that realize these criteria.

- Significant progress on compiler security for safety

Proof techniques

Open questions

How should adversaries be represented? As contexts?

Is there a story for compositional or modular compilation?

We have a class of criteria. Which one(s) should we actually pursue?

What is the connection to other notions of compiler security?


Other notions of compiler security ...

- Gracefully avoiding undefined behavior, *a la* memory-safety (related to our compiler security for safety?)
- Specific defense mechanisms:
 - Stack canaries
 - Code-pointer integrity
 - Address space layout randomization

Is there a uniform, parametric definition of compiler security for these?

- Full abstraction: How does it connect to our notions?
- ...

Conclusion



- Represent external interactions (as traces)
- Represent properties on interactions
- Compiler security = Preservation of a class of properties despite adversaries
- End-to-end security, but secure compilation seems to be broader than this