

Defining Undefined Behavior in Rust

(Work in Progress)

Ralf Jung, Derek Dreyer
Dagstuhl, May 2018

Max Planck Institute for Software Systems (MPI-SWS)



Rust – Mozilla's replacement for C/C++

A safe & modern systems PL



Rust – Mozilla's replacement for C/C++

A safe & **modern** systems PL

- First-class functions
- Polymorphism/generics
- Traits \approx Type classes incl. associated types



Rust – Mozilla's replacement for C/C++

A safe & modern **systems** PL

- First-class functions
- Polymorphism/generics
- Traits \approx Type classes incl. associated types
- Control over resource management (e.g., memory allocation and data layout)



Rust – Mozilla's replacement for C/C++

A **safe** & modern systems PL

- First-class functions
- Polymorphism/generics
- Traits \approx Type classes incl. associated types
- Control over resource management (e.g., memory allocation and data layout)
- Strong type system guarantees:
 - Type & memory safety; absence of data races



Rust – Mozilla's replacement for C/C++



Goal of **RustBelt** project:
Prove safety of Rust and its
standard library.

- First
- Poly
- Trai
- asso
- Con
- (e.g
- Stro

- Type & memory safety; absence of data races

Rust – Mozilla's replacement for C/C++



Goal of **RustBelt** project:
Prove safety of Rust and its
standard library.

That requires defining what Rust is!

- First
- Poly
- Train
- asso
- Con
- (e.g.
- Strong

- Type & memory safety; absence of data races

Aliasing + Mutation



Rust enforces this via **ownership** & **affine types**:



1. **Full ownership**: `T`
 - + *mutation, deallocation*
 - *aliasing*
2. **Mutable reference**: `&mut T`
 - + *mutation*
 - *aliasing, deallocation*
3. **Shared reference**: `&T`
 - + *aliasing*
 - *mutation, deallocation*

Aliasing guarantees

```
fn answer(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique pointers  
    *x = 23;  
    *y = 19;  
    return *x; // must return 23  
}
```

Aliasing guarantees

```
fn answer(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique pointers  
    *x = 23;  
    *y = 19;  
    return *x; // must return 23  
}
```

```
fn shr_read_only(x: &i32) -> bool {  
    // x is read-only  
    let val = *x;  
    unknown_function(x);  
    return *x == val; // must return true  
}
```

Aliasing guarantees

```
fn answer(x: &mut i32, y: &mut i32) -> i32 {
```

Rust's reference types provide
strong aliasing information.

The optimizer should exploit that.

```
    unknown_function(x);
```

```
    return *x == val; // must return true
```

```
}
```

Aliasing guarantees

```
fn answer(x: &mut i32, y: &mut i32) -> i32 {
```

Rust's reference types provide strong aliasing information.

The optimizer should exploit that.

But there is a problem: **unsafe code**.

```
    unknown_function(x);  
    return *x == val; // must return true
```

```
}
```

Unsafe code can access hazardous operations that are banned in safe code:

```
unsafe fn hazardous(x: i32) -> i32 {  
    // *const T is the type of raw (unsafe) pointers  
    let x_ptr = x as *const i32;  
    return *x_ptr; // dereferencing an integer!  
}
```

- Used for better performance, FFI, implementing many standard library types
- Generally encapsulated by safe APIs

```
fn mut_dont_alias(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique pointers  
    *x = 23;  
    *y = 19;  
    return *x; // must return 23  
}
```



```
fn mut_dont_alias(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique pointers  
    *x = 23;  
    *y = 19;  
    return *x; // must return 23  
}
```

Will load 19 because x and y alias!

```
fn main() {  
    let mut x = 0;  
    let x_ptr = &mut x as *mut i32;  
    unsafe { // let's pass the same pointer twice  
        mut_dont_alias(&mut *x_ptr, &mut *x_ptr);  
    }  
}
```

```
fn mut_dont_alias(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique pointers  
    *x = 23;  
    *y = 19;  
    // ...  
}  
  
fn  
1  
1  
unsafe { // let's pass the same pointer twice  
    mut_dont_alias(&mut *x_ptr, &mut *x_ptr);  
}  
}
```

Calling `mut_dont_alias` with aliasing pointers must be UB to justify the optimizations.

Quest for a memory model

The **memory model** for Rust must define how UB arises from unexpected aliasing.

This is work-in-progress. Possible contenders:

1. **Validity-based models** ensure **eagerly** that variables in scope satisfy aliasing restrictions
2. **Access-based models** enforce aliasing restrictions **lazily** when references are used

Validity-based memory model


Every memory location gets equipped with a reader-writer lock.

With $x: \&\text{mut } T$, we hold a write lock.

With $x: \&T$, we hold a read lock.

```
fn mut_dont_alias(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique pointers  
    *x = 23;  
    *y = 19;  
    return *x; // must return 23  
}
```

UB because we attempt to acquire
a write lock the 2nd time



```
fn main() {  
    let mut x = 0;  
    let x_ptr = &mut x as *mut i32;  
    unsafe { // let's pass the same pointer twice  
        mut_dont_alias(&mut *x_ptr, &mut *x_ptr);  
    }  
}
```

```

fn split_at_mut(&mut self, mid: usize) ->
    (&mut [T], &mut [T])
{
    let len = self.len();
    // fn as_mut_ptr(&mut self) -> *mut T
    let ptr = self.as_mut_ptr();
    unsafe {
        assert!(mid <= len);
        // fn from_raw_parts_mut<'a, T>(p: *mut T,
        //                               len: usize) -> &'a mut [T]
        (from_raw_parts_mut(ptr, mid),
         from_raw_parts_mut(ptr.offset(mid as isize),
                             len - mid)
        )
    }
}

```

```
fn split_at_mut(&mut self, mid: usize) ->
    (&mut [T], &mut [T])
{
    let len = self.len();
    // fn as_mut_ptr(&mut self) -> *mut T
    let ptr = self.as_mut_ptr();
    unsafe {
        assert!(mid <= len);
        // fn from_raw_parts_mut<'a, T>(p: *mut T,
        //                                len: usize) -> &'a mut [T]
        (from_raw_parts_mut(ptr, mid),
         from_raw_parts_mut(ptr.offset(mid as isize),
                                id))
    }
}
```

self gets released when
calling as_mut_ptr

UB because locks get acquired on
return value of from_raw_parts_mut,
overlapping with self

Access-based memory model


Keep track of how references are computed, and which reference is *active* for a location.

On every memory access:

- The active reference must be “derived from” the accessing reference.
- The accessing reference becomes the new active reference.

Creating a new reference performs the same check, then makes the new reference active.


```
fn mut_dont_alias(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique pointers  
    *x = 23;  
    *y = 19;  
    return *x; // must return 23  
}
```



UB because x (last reference to access,
and hence **active**) is not “derived from” y.

```
fn main() {  
    let mut x = 0;  
    let x_ptr = &mut x as *mut i32;  
    unsafe { // let's pass the same pointer twice  
        mut_dont_alias(&mut *x_ptr, &mut *x_ptr);  
    }  
}
```

```

fn split_at_mut(&mut self, mid: usize) ->
    (&mut [T], &mut [T])
{
    let len = self.len();
    // fn as_mut_ptr(&mut self) -> *mut T
    let ptr = self.as_mut_ptr();
    unsafe {
        assert!(mid <= len);
        // fn from_raw_parts_mut<'a, T>(p: *mut T,
        //                               len: usize) -> &'a mut [T]
        (from_raw_parts_mut(ptr, mid),
         from_raw_parts_mut(ptr.offset(mid as isize),
                             len - mid)
        )
    }
}

```

```
fn split_at_mut(&mut self, mid: usize) ->
```

```
    (&mut [T], &mut [T])
```

```
{
```

```
    let len = self.len();
```

```
    // fn as_mut_ptr(&mut self) -> *mut T
```

```
    let ptr = self.as_mut_ptr();
```

```
    unsafe {
```

```
        assert!(mid <= len);
```

```
        // fn from_raw_parts_mut<'a, T>(p: *mut T,
```

```
        //                               len: usize) -> &'a mut [T]
```

```
        (from_raw_parts_mut(ptr, mid),
```

```
         from_raw_parts_mut(ptr.offset(mid as isize),
```

```
         mid)
```

Return value now is active
reference for first part, no conflict.

```
}
```

`self` is not ever used again, hence
never becomes active again.

Access-based breaks some optimizations

Suppose `x`: `&i32` is in scope.

```
loop {  
    ...  
    let y = *x;  
    ...  
}  
  
    ~~~~  
  
let tmp = *x;  
loop {  
    ...  
    let y = tmp;  
    ...  
}
```

Access-based breaks some optimizations

Suppose `x`: `&i32` is in scope.

```
loop {  
    ...  
    let y = *x;  
    ...  
}  
  
let tmp = *x;  
loop {  
    ...  
    let y = tmp;  
    ...  
}
```

Problem: What if `*x` is UB, but the code running before it breaks out of the loop?

Memory model trade-offs

Validity-based model:

- Complex to define and reason about
- Program behavior can depend on “lifetime” of reference as inferred by the compiler
- Difficult to make compatible with existing unsafe code

Access-based model:

- Catches errors very late, makes it hard to assign blame
- Requires tracking pointer provenance
- Allows more code, and hence fewer optimizations