



Is the CakeML compiler secure?

Dagstuhl Seminar on Secure Compilation, May 2018

Active developers: Oskar Abrahamsson, Johannes Åman Pohjola, Hugo Férée, Anthony Fox, **Ramana Kumar**, Andreas Lööw, Alexander Mihajlovic, **Magnus Myreen, Michael Norrish, Scott Owens, Yong Kiam Tan**





Is the CakeML compiler secure?

This talk:

1. What is the CakeML compiler?
2. What has been proved about it? Is it secure?
3. Discussion: how could it be made more secure?

What is the CakeML compiler?

Aim: To be a **realistic** ML compiler that can be used for **research and teaching** (e.g. MSc students).

Realistic end-points:

no undefined behaviour

High-level ML source language that is convenient for program verification: **no bounds on sizes** and **mathematical integers** as default int type; **type safe**.

Output: **concrete machine code** (not assembly) with semantics that is as accurate as possible for **commercial ISAs** (but no concurrency).

Acceptable performance, bootstrapped compiler etc.

Implementation

Latest version:

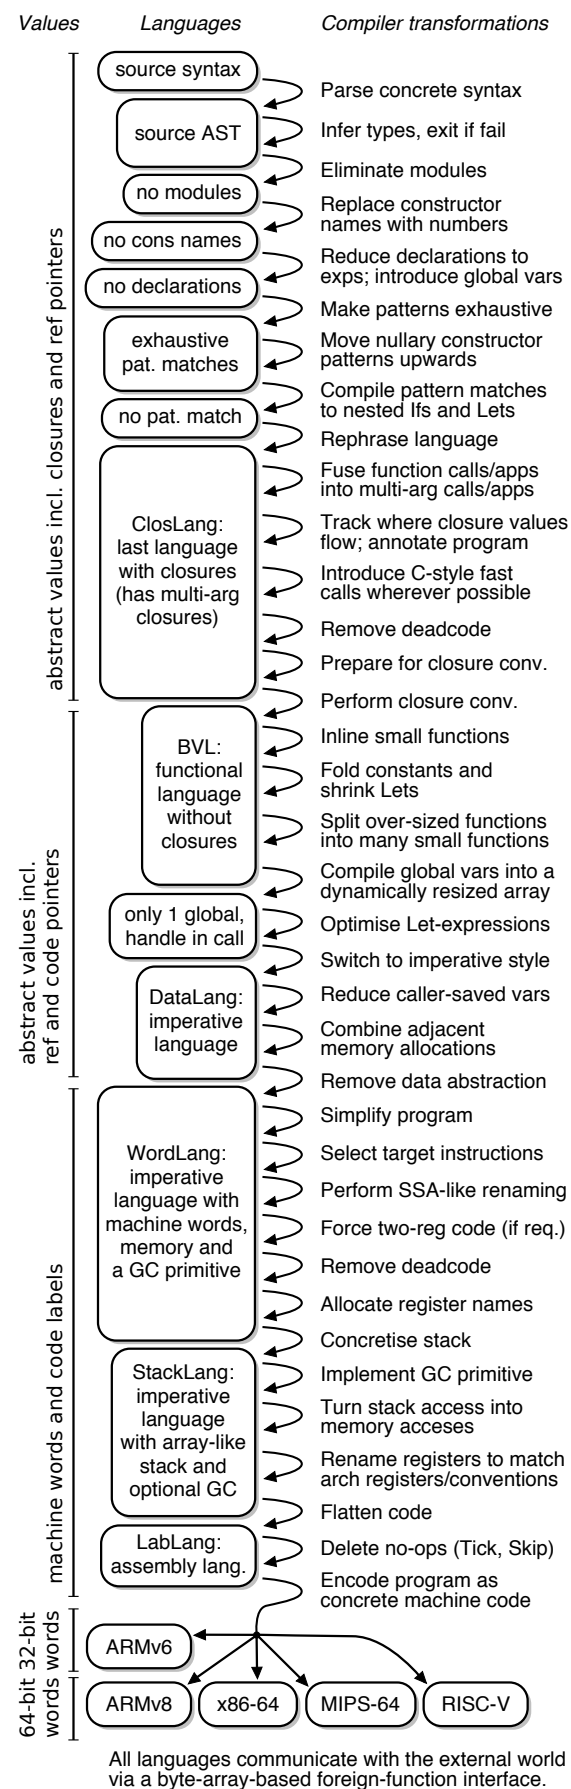
12 intermediate languages (ILs)

and many within-IL optimisations

each IL at the right level of abstraction

for the benefit of
proofs and compiler
implementation

Next slide zooms in



Values used by
the semantics

Both proved sound
and complete.

Values

Languages

Compiler transformations

source syntax

Parse concrete syntax

source AST

Infer types, exit if fail

no modules

Eliminate modules

no cons names

Replace constructor
names with numbers

no declarations

Reduce declarations to
exps; introduce global vars

exhaustive
pat. matches

Make patterns exhaustive

Move nullary constructor
patterns upwards

no pat. match

Compile pattern matches
to nested ifs and Lets

Rephrase language

ClosLang:
last language
with closures
(has multi-arg
closures)

Fuse function calls/apps
into multi-arg calls/apps

Track where closure values
flow; annotate program

Introduce C-style fast
calls wherever possible

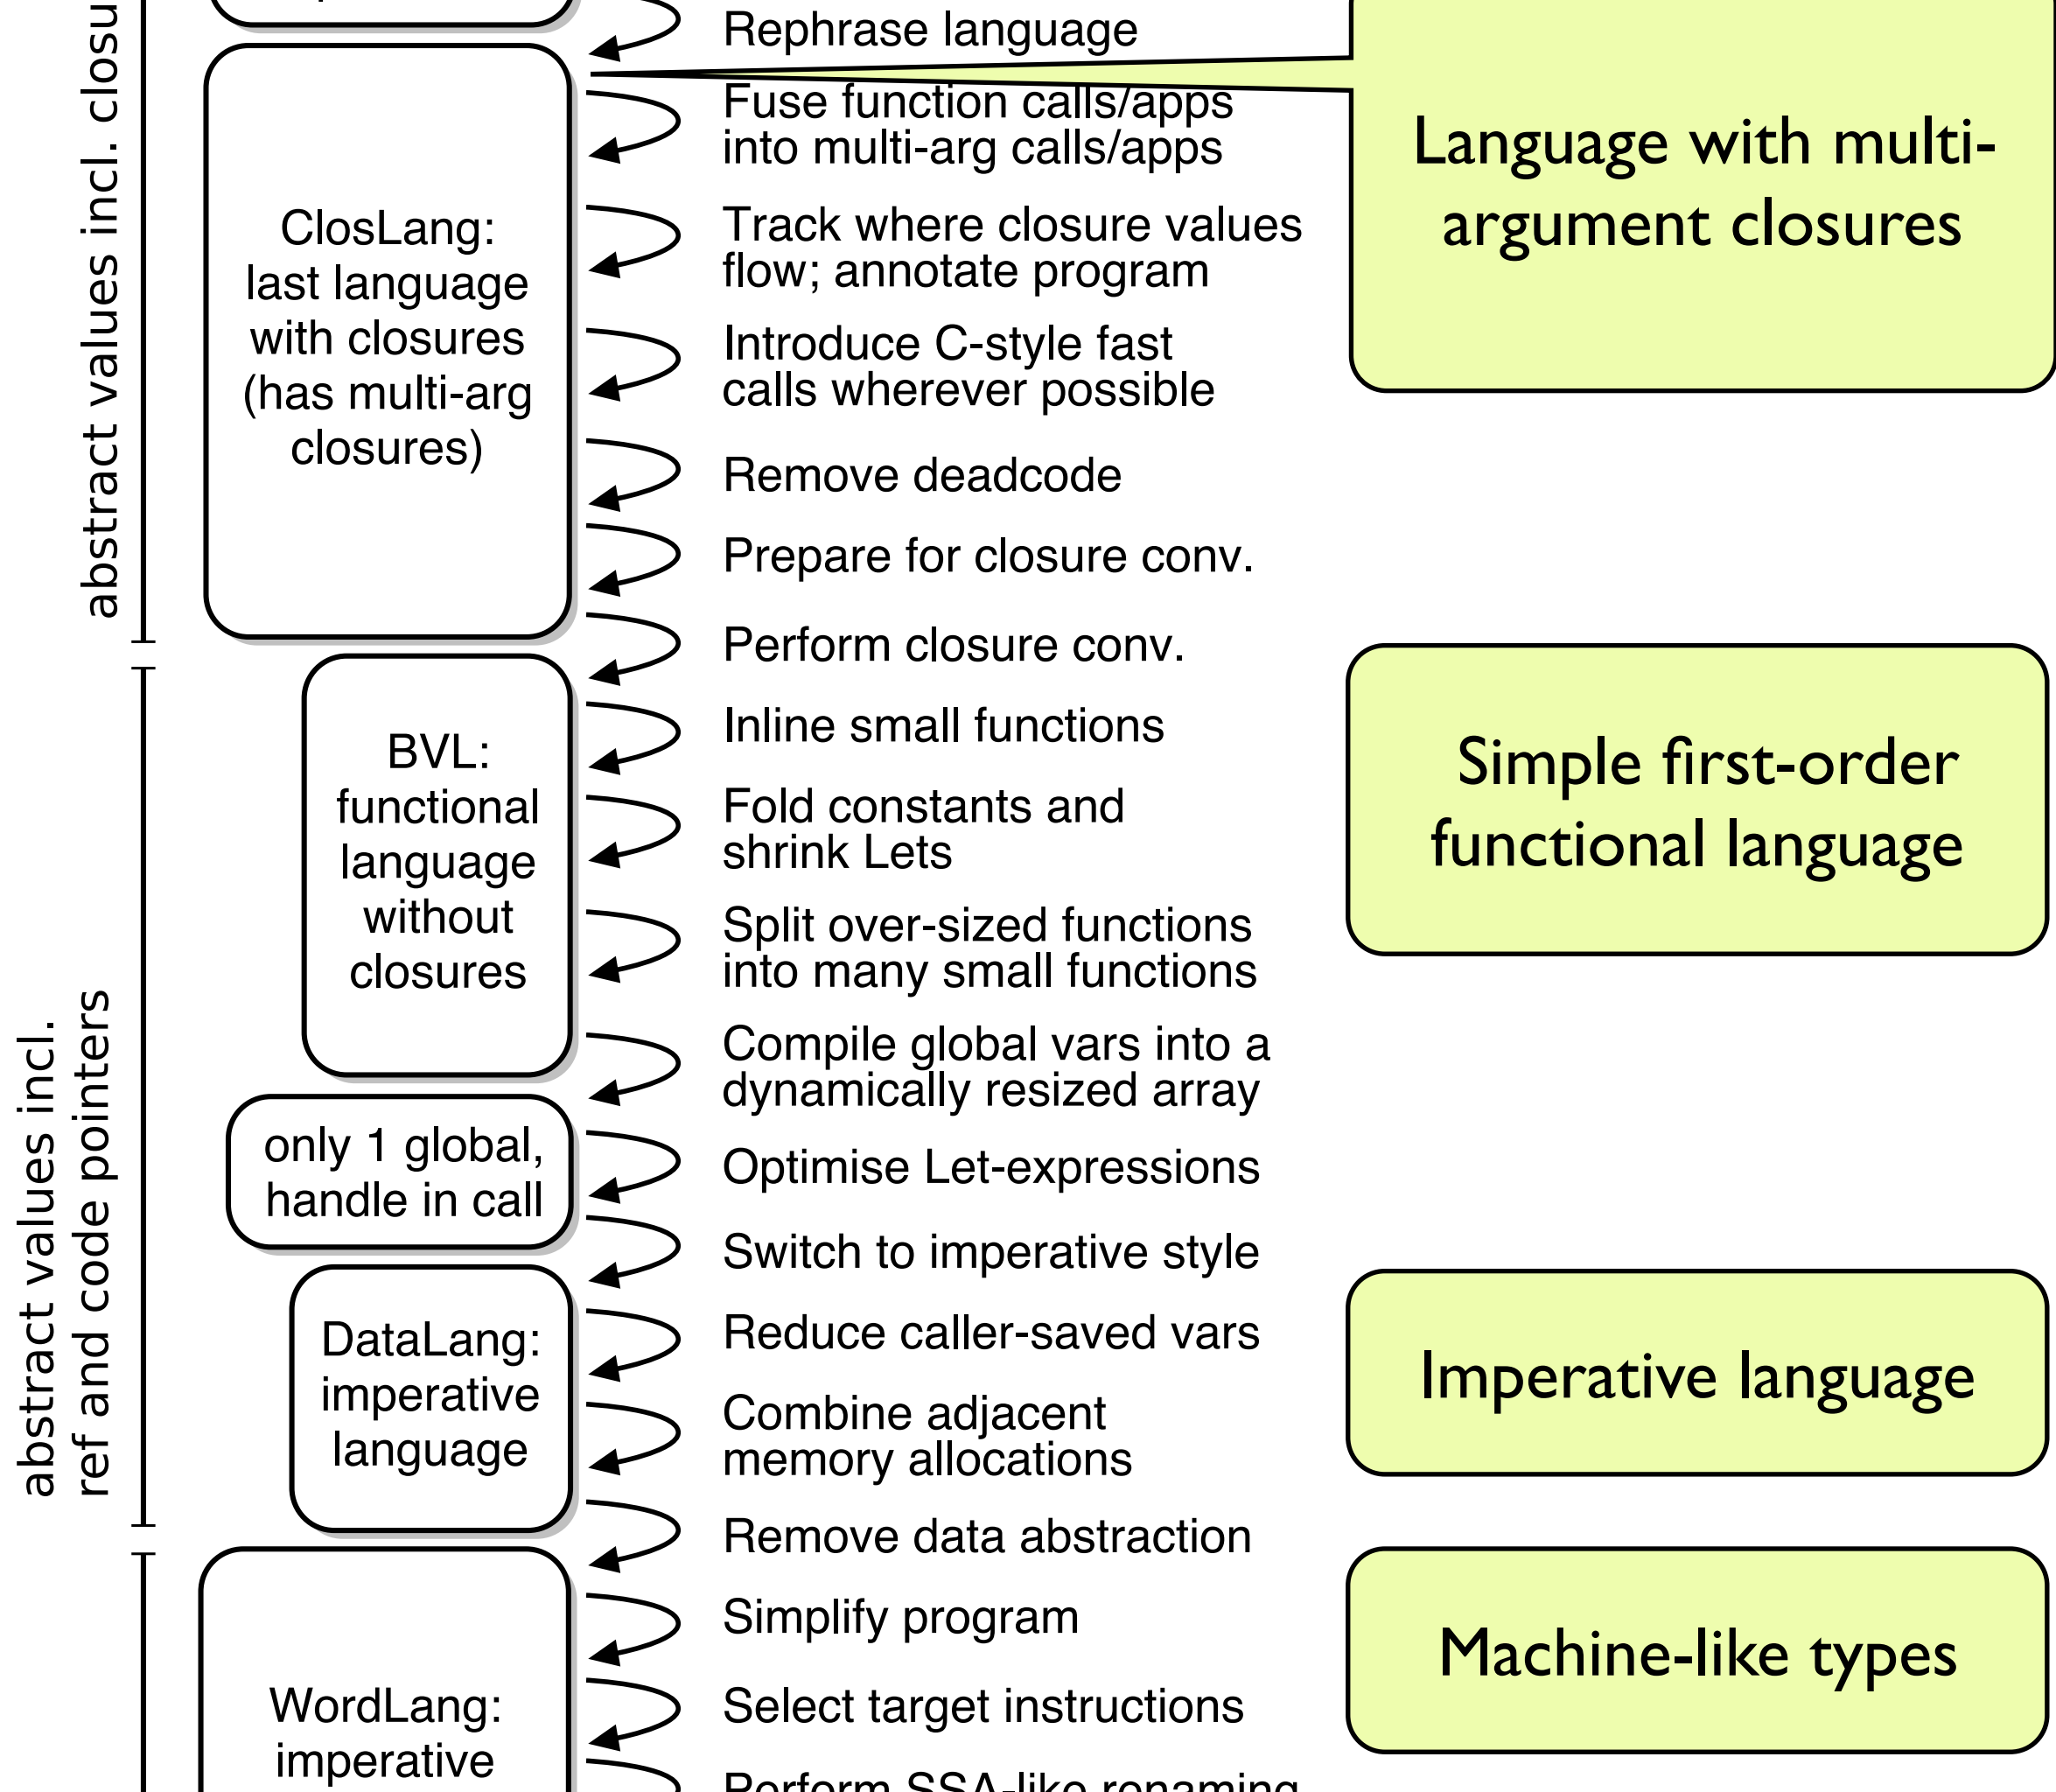
Remove deadcode

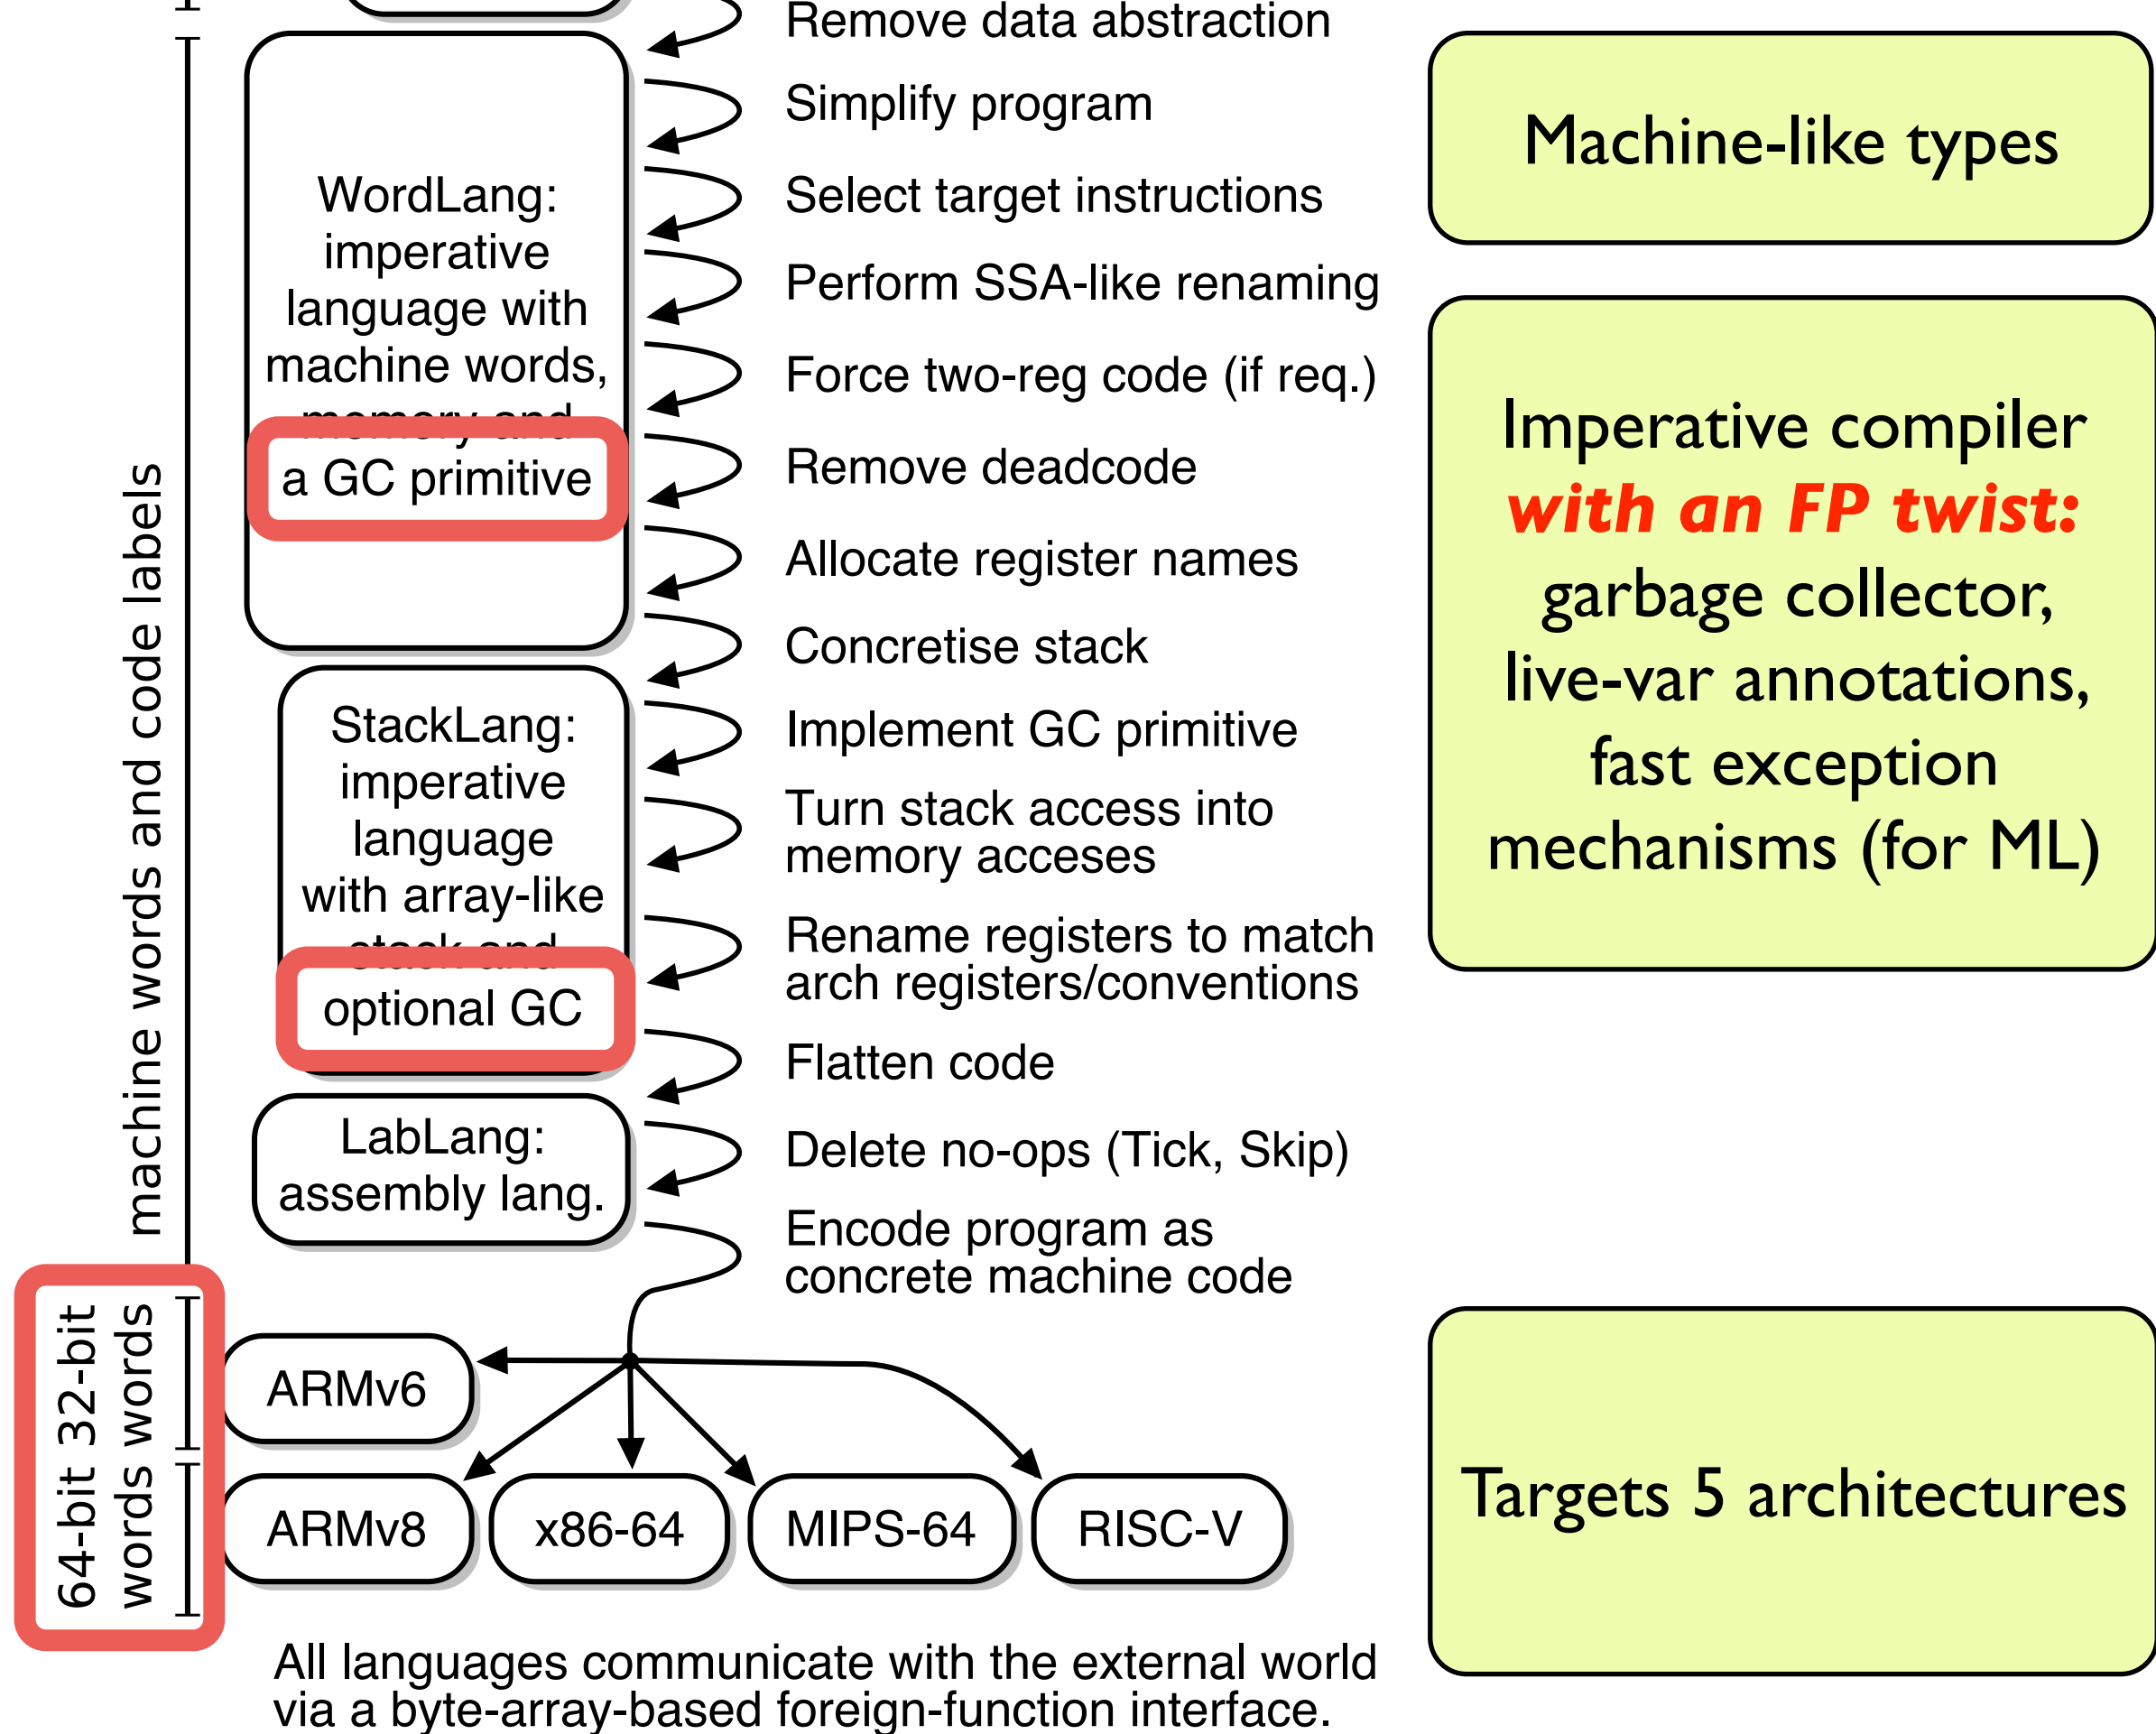
Parser and type
inferencer as before

Early phases reduce
the number of
language features

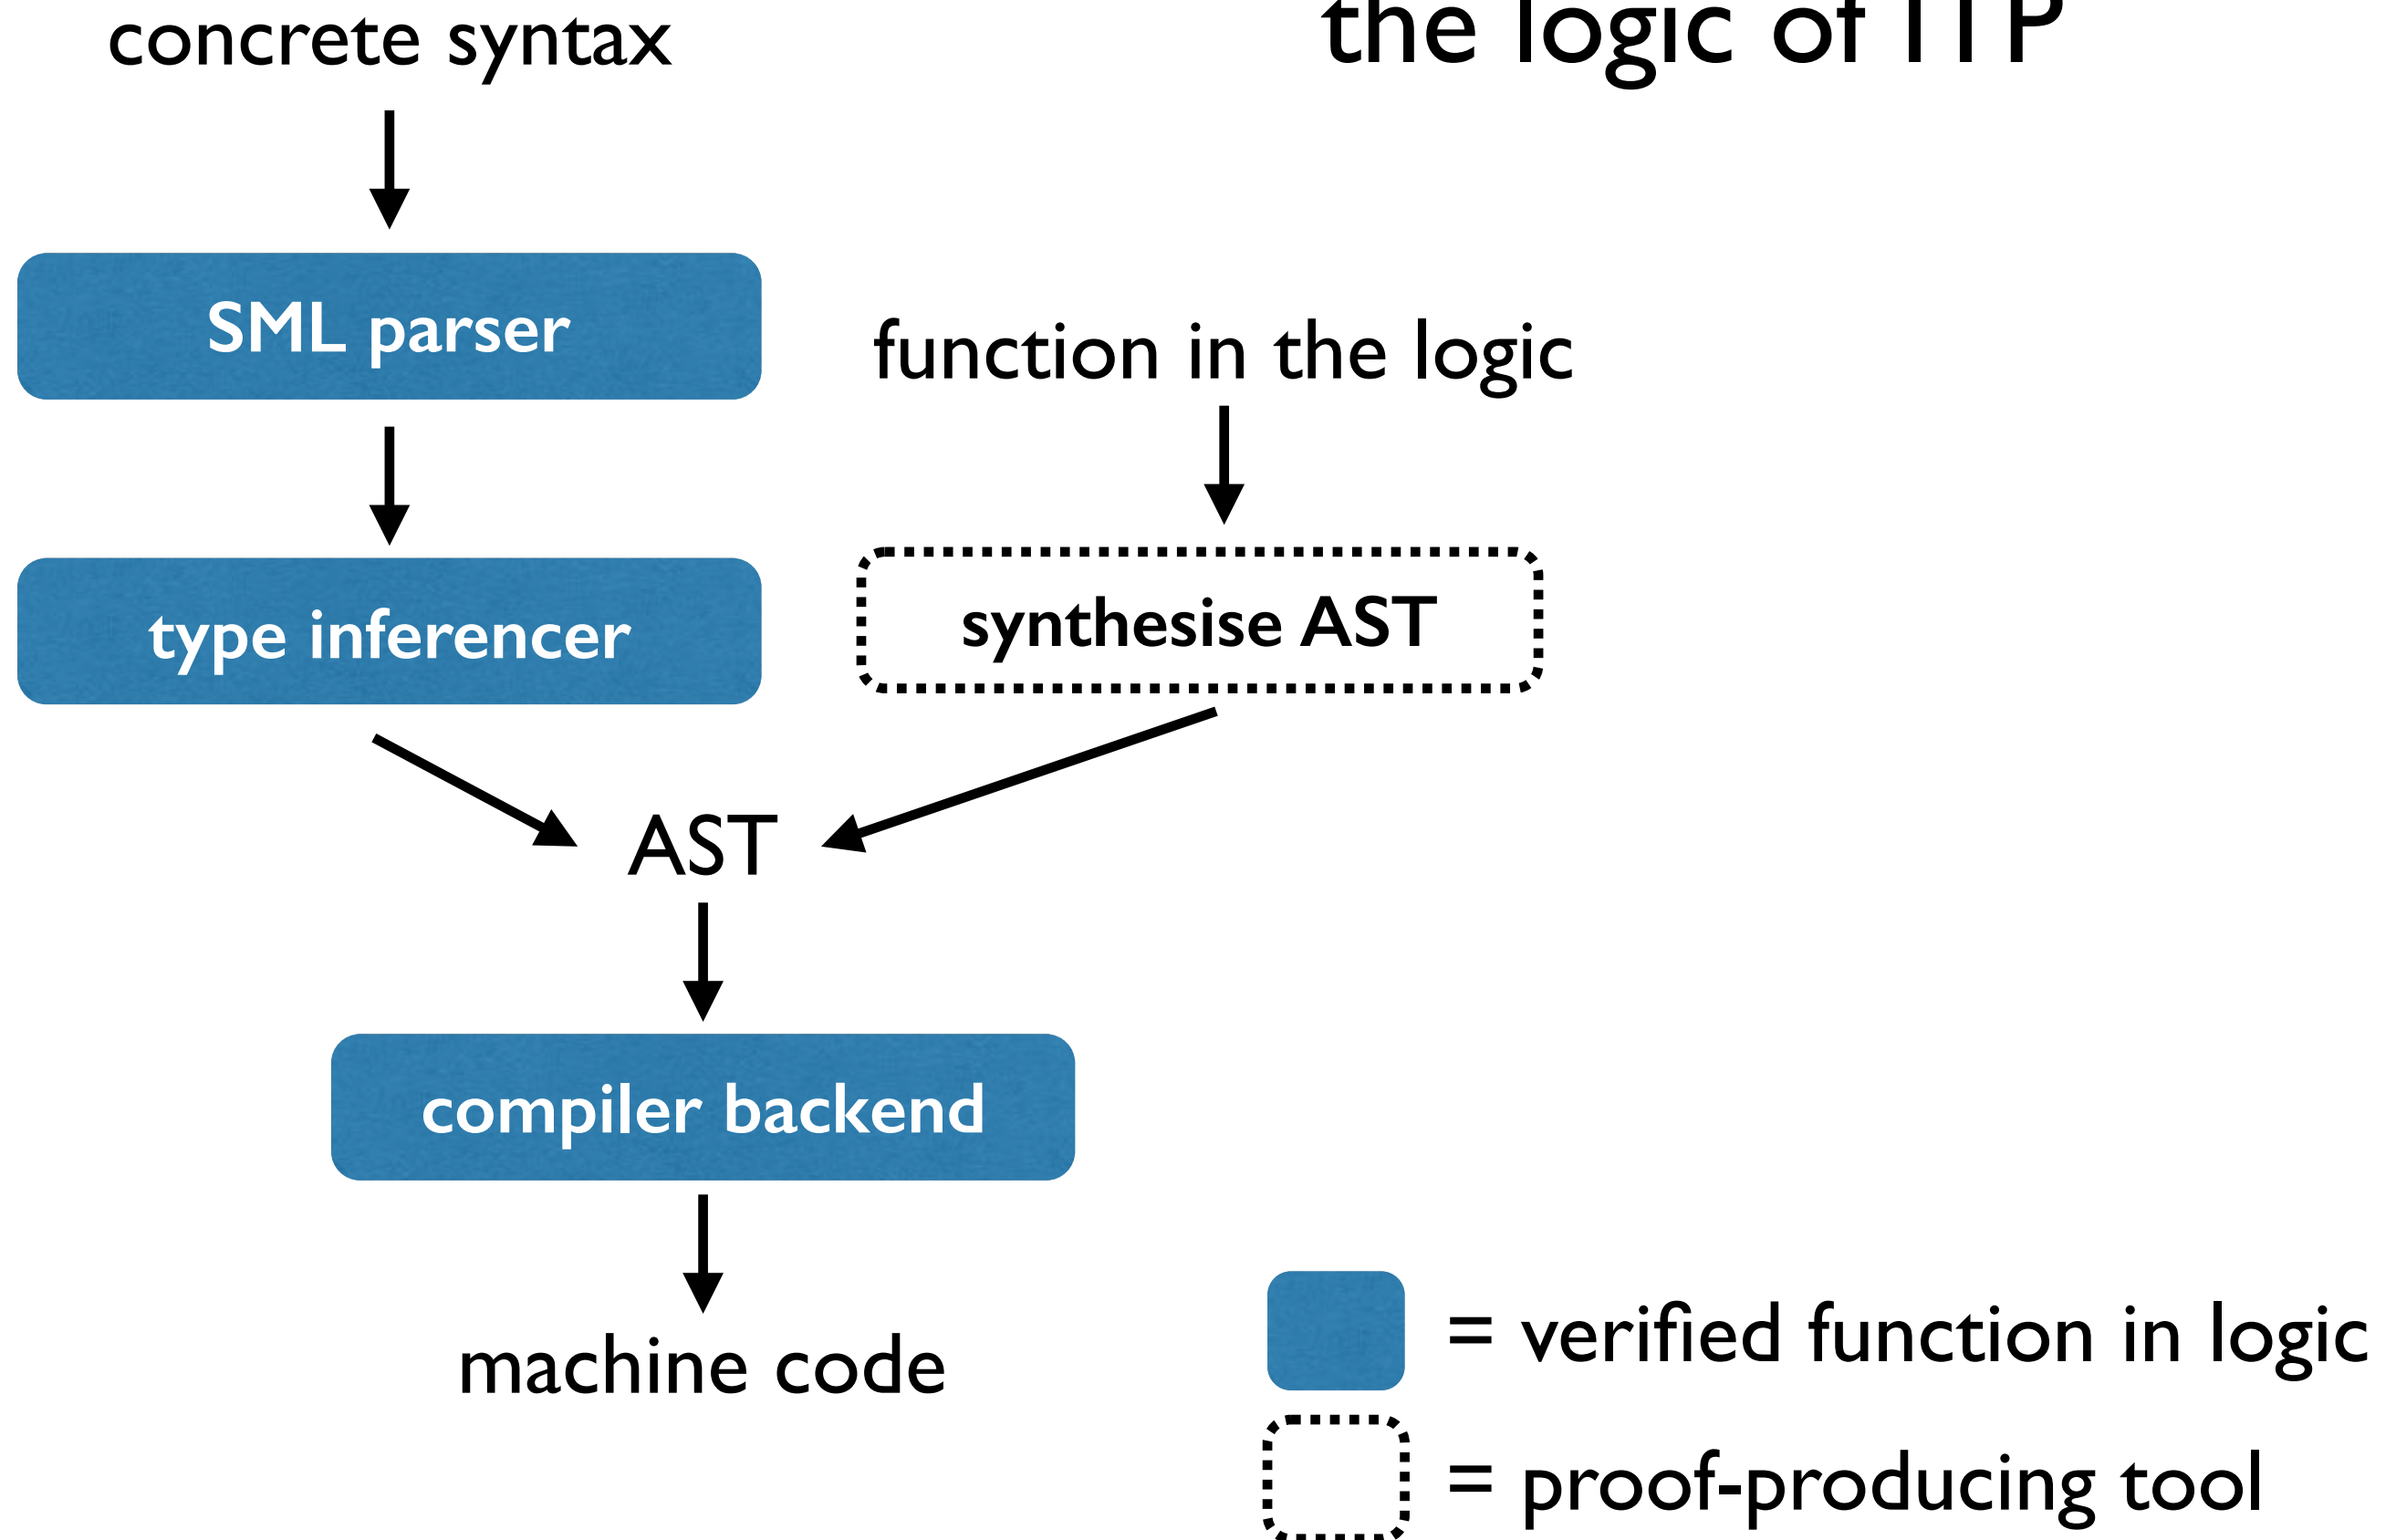
Language with multi-
argument closures

tract values incl. closures and ref pointers





Bootstrapping inside the logic of ITP




input: parser + type inferencer + compiler backend

function in the logic



synthesise AST



AST



compiler backend




machine code

output: verified implementation of
parser + type inf. + backend




= verified function in logic



= proof-producing tool

Part 2:

This talk:

1. What is the CakeML compiler?
-  2. What has been proved about it? Is it secure?
3. Discussion: how could it be made more secure?

Compilation cannot be perfect

Source code:

- no bounds on length of lists, trees etc.
- mathematical integers

bignum lib
& GC



*source programs cannot always be flawlessly
represented in the target languages*

Target languages:

- 32-bit and 64-bit architectures
- finite memory

Top-level correctness theorem

(next slides will explain more)

$\vdash \text{inf_conf_ok } cc.\text{inferencer_config} \wedge \neg cc.\text{input_is_sexp} \wedge$
 $\text{backend_config_ok } cc.\text{backend_config} \wedge \text{mc_conf_ok } mc \wedge$
 $\text{mc_init_ok } cc.\text{backend_config } mc \Rightarrow$
case $\text{cakeml_compile } cc \text{ prelude input}$ of
 $\text{Success } (code, data, cc') \Rightarrow$
 $\exists \text{behaviours.}$
 $\text{cakeml_semantics } ffi \text{ prelude input} = \text{Execute } behaviours \wedge$
 $\forall ms.$
 $\text{installed } code \text{ data } cc'.ffi_names \text{ ffi}$
 $(\text{heap_regs } cc.\text{backend_config}.\text{stack_conf}.\text{reg_names}) \text{ mc } ms \Rightarrow$
 $\text{machine_sem } mc \text{ ffi } ms \subseteq \text{extend_with_resource_limit } behaviours$
| $\text{Failure ParseError} \Rightarrow \text{cakeml_semantics } ffi \text{ prelude input} = \text{CannotParse}$
| $\text{Failure (TypeError _)} \Rightarrow \text{cakeml_semantics } ffi \text{ prelude input} = \text{IllTyped}$
| $\text{Failure CompileError} \Rightarrow \text{true}$
| $\text{Failure (ConfigError _)} \Rightarrow \text{true}$

Top-level correctness theorem

$\vdash \text{inf_conf_ok } cc.\text{inferencer_config} \wedge \neg cc.\text{input_ok} \wedge$
 $\text{backend_config_ok } cc.\text{backend_config} \wedge mc$
 $\text{mc_init_ok } cc.\text{backend_config } mc \Rightarrow$
 $\text{case } \text{cakeml_compile } cc \text{ prelude input of}$
 $\text{Success } (code, data, cc') \Rightarrow$
 $\exists \text{ behaviours.}$
 $\text{cakeml_semantics } ffi \text{ prelude input} = \text{Execute } behaviours \wedge$
 $\forall ms.$
 $\text{installed } code \text{ data } cc'.ffi_names \text{ ffi}$
 $(\text{heap_regs } cc.\text{backend_config}.\text{stack_conf}.\text{reg_names}) \text{ mc } ms \Rightarrow$
 $\text{machine_sem } mc \text{ ffi } ms \subseteq \text{extend_with_resource_limit } behaviours$
| $\text{Failure ParseError} \Rightarrow \text{cakeml_semantics } ffi \text{ prelude input} = \text{CannotParse}$
| $\text{Failure (TypeError _)} \Rightarrow \text{cakeml_semantics } ffi \text{ prelude input} = \text{IllTyped}$
| $\text{Failure CompileError} \Rightarrow \text{true}$
| $\text{Failure (ConfigError _)} \Rightarrow \text{true}$

if the compiler returns
success then the source code
has well-defined behaviour

Error cases

$\vdash \text{inf_conf_ok } cc.\text{inferencer_config} \wedge \neg cc.\text{input_is_sexp} \wedge$
 $\text{backend_config_ok } cc.\text{backend_config} \wedge \text{mc_conf_ok } mc \wedge$
 $\text{mc_init_ok } cc.\text{backend_config } mc \Rightarrow$
case `cakeml_compile cc prelude input` of
 Success $(code, data, cc')$ \Rightarrow
 $\exists \text{behaviours}.$
 $\text{cakeml_semantics } ffi \text{ prelude input} = \text{Execute } behaviours \wedge$
 $\forall ms.$
 installed $code \ data \ cc'.ffi_names \ ffi$
 $(\text{heap_regs } cc.\text{backend_config}.\text{stack_conf}.\text{reg_names}) \ mc \ ms \Rightarrow$
 $\text{machine_sem } mc \ ffi \ ms \subseteq \text{extend_with_resource_limit } behaviours$
| Failure ParseError $\Rightarrow \text{cakeml_semantics } ffi \text{ prelude input} = \text{CannotParse}$
| Failure (TypeError _) $\Rightarrow \text{cakeml_semantics } ffi \text{ prelude input} = \text{IllTyped}$
| Failure CompileError $\Rightarrow \text{true}$
| Failure (ConfigError _) $\Rightarrow \text{true}$

Partiality I (at compile time)

$\vdash \text{inf_conf_ok } cc.\text{inferencer_config} \wedge \neg cc.\text{input_is_sexp} \wedge$
 $\text{backend_config_ok } cc.\text{backend_config} \wedge \text{mc_conf_ok } mc \wedge$
 $\text{mc_init_ok } cc.\text{backend_config } mc \Rightarrow$
 $\text{case cakeml_compile } cc \text{ prelude input of}$
 $\text{Success } (code, data, cc') \Rightarrow$
 $\exists \text{ behaviours.}$
 $\text{cakeml_semantics ffi prelude input} = \text{Execute behaviours} \wedge$
 $\forall ms.$
 $\text{installed code data } cc'.\text{ffi_names ffi}$
 $(\text{heap_regs } cc.\text{backend_config.stack_conf.reg_names}) mc ms \Rightarrow$
 $\text{machine_sem } mc \text{ ffi } ms \subseteq \text{extend_with_resource_limit behaviours}$
| $\text{Failure ParseError} \Rightarrow \text{cakeml_semantics ffi prelude input} = \text{CannotParse}$
| $\text{Failure (TypeError _)} \Rightarrow \text{cakeml_semantics ffi prelude input} = \text{IllTyped}$
| $\text{Failure CompileError} \Rightarrow \text{true}$
| $\text{Failure (ConfigError _)} \Rightarrow \text{true}$

compiler can exit with
error (rare in practice)

Partiality 2 (at execution time)

$\vdash \text{inf_conf_ok } cc.\text{inferencer_config} \wedge \neg cc.\text{input_is_sexp} \wedge$
 $\text{backend_config_ok } cc.\text{backend_config} \wedge \text{mc_conf_ok } mc \wedge$
 $\text{mc_init_ok } cc.\text{backend_config } mc \Rightarrow$

$\text{case cakeml_compile } cc \text{ prelude input of}$
 $\text{Success } (code, data, cc') \Rightarrow$
 $\exists \text{ behaviours.}$

$\text{cakeml_semantics ffi prelude input} = \text{Execute behaviours} \wedge$
 $\forall ms.$

$\text{installed code data } cc'.\text{ffi_names ffi}$

$(\text{heap_regs } cc.\text{backend_config.stack_conf.reg_names}) mc ms \Rightarrow$

$\text{machine_sem } mc \text{ ffi } ms \subseteq \text{extend_with_resource_limit behaviours}$

- | $\text{Failure ParseError} \Rightarrow \text{cakeml_semantics ffi prelude input} = \text{CannotParse}$
- | $\text{Failure (TypeError _)} \Rightarrow \text{cakeml_semantics ffi prelude input} = \text{IllTyped}$
- | $\text{Failure CompileError} \Rightarrow \text{true}$
- | $\text{Failure (ConfigError _)} \Rightarrow \text{true}$

executable can give up with
out-of-memory (OOM) error

Information leakage?

Behaviour can differ

\exists *behaviours*.

cakeml_semantics ffi prelude input = Execute *behaviours* \wedge

...

machine_sem mc ffi ms \subseteq extend_with_resource_limit *behaviours*

under cost limit implies no OOM

One could have a **cost semantics**

... but this would clutter the source semantics

How would one reason about the cost semantics? (not local)

Would we need to preserve OOM?

Timing

Nothing is proved about timing.

Constant-time code impossible to write in CakeML.

Reasons:

- GC,
- bignums,
- function call optimisations,
- unpredictable timing in target ISAs

Assumptions about exec. env.

$\vdash \text{inf_conf_ok } cc.\text{inferenc}$
 $\text{backend_config_ok } cc$
 $\text{mc_init_ok } cc.\text{backend_config}$
 $\text{case cakeml_compile } cc$
 $\text{Success } (code, data, cc')$
 $\exists \text{behaviours}$
 $\text{cakeml_semantics } ffi$
 $\forall ms.$

we assume that the machine code is present in the machine state

... and that each foreign function interface (FFI) call only touches FFI relevant state and behaves according to *ffi* oracle model.

installed *code data cc'.ffi_names ffi*

(heap_regs *cc.backend_config.stack_conf.reg_names*) *mc ms* \Rightarrow

machine_sem mc ffi ms $\subseteq \text{extend_with_resource_limit behaviours}$

| Failure ParseError $\Rightarrow \text{cakeml_semantics } ffi \text{ prelude input} = \text{CannotParse}$

| Failure (T
 | Failure C
 | Failure (C
 we allow the operating system and other processes run
 between each instruction execution *and assume that*
other processes do not alter CakeML's memory

Reasonable assumptions?


What are reasonable assumptions about the execution environment?

Answer(?): assumptions that could be discharged if the proof is composed with verified OS, verified linker etc.

... but where do we stop?

Part 3:

This talk:

1. What is the CakeML compiler?
2. What has been proved about it? Is it secure?
-  3. Discussion: how could it be made more secure?

Summary

CakeML is safe

safe control-flow (no return-oriented programming), no buffer overflows, ...

but we could have
a cost model

CakeML is only secure to weak attacker model

attacker needs to be blind to OOM

attacker has very coarse sense of time

realistic only for
some applications

attacker only has access through CakeML's FFI

reasonable assumption?